

```

import numpy as np
import pandas as pd
import torch
from torch import optim
import torch.nn as nn
from PIL import Image
import os

from skimage import filters, color, morphology, io
import matplotlib.pyplot as plt
import numpy as np
import random
import skimage as si

from google.colab import auth
from google.colab import drive

import gspread

from google.auth import default
from PIL import Image

#authenticating to google
auth.authenticate_user()
creds, _ = default()
gc = gspread.authorize(creds)

drive.mount('/content/drive')

    Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

sh = gc.open_by_url('https://docs.google.com/spreadsheets/d/1eNgK86cm6L10VI5iavk7mY3UQUd1VHeXgwXSTyNNuRw/edit?usp=sharing')
ws = sh.worksheet('train')

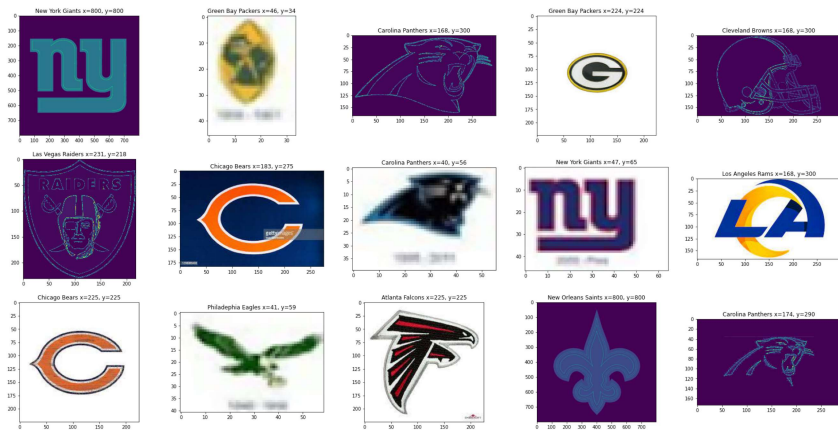
df = pd.DataFrame(ws.get_all_records())

n1 = 3
n2 = 5

fig, axs = plt.subplots(n1, n2, figsize=(32, 16))

for i in range(n1):
    for j in range(n2):
        z = random.randint(0, len(df)-1)
        pic = np.array(Image.open(df.loc[z]["filepath"]))
        shape = np.shape(pic)
        axs[i, j].imshow(pic)
        axs[i, j].set_title('%s x=%.f, y=%.f' % (df.loc[z]["team_name"], shape[0], shape[1]))

```



```
import numpy as np
```

```
sample_df = df.sample(frac=0.01, ignore_index=True)
```

```
N = len(sample_df)
shape = np.zeros((2, N))
```

```
for i in range(N):
    tmp = np.shape(np.array(Image.open(sample_df.loc[i]["filepath"]).convert('L'))))
    shape[:, i] = [tmp[0], tmp[1]]
```

```
fig, axs = plt.subplots(1, 2, figsize=(20, 5))
```

```
axs[0].scatter(shape[0, :], shape[1, :])
axs[0].plot(range(0, 1000), range(0, 1000), 'k')
```

```


axs[0].set_ylabel('dim 1')
axs[0].set_xlabel('dim 0')
axs[0].grid()
axs[0].set_title('Scatter of images shapes')
axs[0].set_xlim([0, 1000])
axs[0].set_ylim([0, 1000])

axs[1].hist(shape[0, :]/shape[1, :])
axs[1].grid()
axs[1].set_xlabel('Aspect Ratio')
axs[1].set_ylabel('Frequency')
axs[1].set_title('Histogram of aspect ration for images')
axs[1].set_xlim([0, 2])

print("Average height " + str(sum(shape[1, :]) / len(shape[1, :])))

```

Average height 161.0



```

from torch.utils.data import Dataset

class CustomDataset(Dataset):
    def __init__(self, X, y, BatchSize, transform):
        super().__init__()
        self.BatchSize = BatchSize
        self.y = y
        self.X = X
        self.transform = transform

    def num_of_batches(self):
        """
        Detect the total number of batches
        """
        return math.floor(len(self.list_IDs) / self.BatchSize)

    def __getitem__(self, idx):
        class_id = self.y[idx]
        img = Image.open(self.X[idx])
        img = img.convert("RGBA").convert("RGB")
        img = self.transform(img)
        return img, torch.tensor(int(class_id))

```

```

def __len__(self):
    return len(self.X)

from sklearn.model_selection import train_test_split
from torch.utils.data import DataLoader
from torchvision import transforms

# Shuffle dataframe
df = df.sample(frac=1)

X = df.iloc[:,0]
y = df.iloc[:,2]

transform = transforms.Compose([
    transforms.Resize([256,256]),
    transforms.RandomRotation(20, fill=256),
    transforms.ToTensor(),
    transforms.RandomAffine(degrees=0, translate=(0.025, 0.025), fill=256),
    transforms.Normalize([0.5], [0.5])
])

test_transform = transforms.Compose([
    transforms.Resize([256,256]),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)),
])

train_ratio = 0.80
validation_ratio = 0.1
test_ratio = 0.1
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1 - train_ratio, stratify = y, random_state = 0)
X_val, X_test, y_val, y_test = train_test_split(X_test, y_test, test_size=test_ratio/(test_ratio + validation_ratio), random_state = 0)

dataset_stages = ['train', 'val', 'test']

batch_size = 32
image_datasets = {'train' : CustomDataset(X_train.values, y_train.values, batch_size, transform), 'val' : CustomDataset(X_val.values, y_val.values, batch_size, test_transform), 'test' : Cu
dataloaders = {x: DataLoader(image_datasets[x], batch_size=image_datasets[x].BatchSize,
                             shuffle=True, num_workers=0)
               for x in dataset_stages}

dataset_sizes = {x: len(image_datasets[x]) for x in dataset_stages}

print(dataset_sizes)

{'train': 256, 'val': 32, 'test': 32}

nparray = image_datasets['train'][12][0].cpu().numpy()
image = transforms.ToPILImage()(image_datasets['train'][12][0].cpu()).convert("RGB")
display(image)

```



```
import time

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

def train_model(model, criterion, optimizer, scheduler, num_epochs=15):
    since = time.time()
    best_acc = 0.0

    for epoch in range(num_epochs):
        print('Epoch {}/{}'.format(epoch, num_epochs - 1))
        print('-' * 10)
        # Each epoch has a training and validation phase
        for phase in ['train', 'val']:
            if phase == 'train':
                model.train() # Set model to training mode
            else:
                model.eval() # Set model to evaluate mode

        running_loss = 0.0
        running_corrects = 0
        num_batches = 0
        outputs = None
        # Iterate over data.
        for inputs, labels in dataloaders[phase]:
            # Loading Bar
            if (phase == 'train'):
                num_batches += 1
                percentage_complete = ((num_batches * batch_size) / (dataset_sizes[phase])) * 100
                percentage_complete = np.clip(percentage_complete, 0, 100)
                print("{:0.2f}".format(percentage_complete), "% complete", end="\r")

            inputs = inputs.to(device)
            labels = labels.to(device)

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward
            # track history if only in train
            with torch.set_grad_enabled(phase == 'train'):
                outputs = model(inputs)
                loss = criterion(outputs.float(), labels)

            # backward + optimize only if in training phase
            if phase == 'train':
                loss.backward()
                # TODO: try removal
                torch.nn.utils.clip_grad_norm_(model.parameters(), 1)
```

```

optimizer.step()

# statistics
running_loss += loss.item() * inputs.size(0)

predicted = torch.max(outputs.data, 1)[1]
running_correct = (predicted == labels).sum()
running_corrects += running_correct
if phase == 'train':
    scheduler.step()

epoch_loss = running_loss / dataset_sizes[phase]

epoch_acc = running_corrects / dataset_sizes[phase]
#epoch_acc = sum(epoch_acc) / len(epoch_acc)

print('{} Loss: {:.4f} Acc: {:.4f}'.format(
    phase, epoch_loss, epoch_acc.item()))

time_elapsed = time.time() - since
print('Training complete in {:.0f}m {:.0f}s'.format(
    time_elapsed // 60, time_elapsed % 60))
return model

from torchvision import models
from torch.optim import lr_scheduler

model_ft = models.squeezenet1_1(pretrained=True)
model_ft.num_classes = 32 #this is an important component of the output model
model_ft.classifier._modules["1"] = nn.Conv2d(512, model_ft.num_classes, kernel_size=(1, 1))
for param in model_ft.parameters():
    param.requires_grad = False
for param in model_ft.classifier.parameters():
    param.requires_grad = True

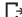
criterion = nn.CrossEntropyLoss()

optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.01)

exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)

model_ft = train_model(model_ft.to(device), criterion, optimizer_ft, exp_lr_scheduler, 15)

```

 Epoch 1/14

 train Loss: 195.6700 Acc: 0.0312
 val Loss: 8.3117 Acc: 0.2188
 Epoch 2/14

 train Loss: 128.1705 Acc: 0.0352
 val Loss: 6.5645 Acc: 0.2500

```

Epoch 5/14
-----
train Loss: 16.0656 Acc: 0.1250
val Loss: 3.7918 Acc: 0.3750
Epoch 6/14
-----
train Loss: 3.4933 Acc: 0.4219
val Loss: 2.6368 Acc: 0.4688
Epoch 7/14
-----
train Loss: 1.8375 Acc: 0.5977
val Loss: 2.3684 Acc: 0.5312
Epoch 8/14
-----
train Loss: 1.8471 Acc: 0.6094
val Loss: 2.0991 Acc: 0.5312
Epoch 9/14
-----
train Loss: 1.6187 Acc: 0.6328
val Loss: 1.9921 Acc: 0.5000
Epoch 10/14
-----
train Loss: 1.5550 Acc: 0.6484
val Loss: 1.9330 Acc: 0.5000
Epoch 11/14
-----
train Loss: 1.3840 Acc: 0.6797
val Loss: 1.8805 Acc: 0.5312
Epoch 12/14
-----
train Loss: 1.4538 Acc: 0.6797
val Loss: 1.8304 Acc: 0.5312
Epoch 13/14
-----
train Loss: 1.3677 Acc: 0.7148
val Loss: 1.7130 Acc: 0.5625
Epoch 14/14
-----
train Loss: 1.2909 Acc: 0.7578
val Loss: 1.7002 Acc: 0.5312

```

```
from sklearn.metrics import accuracy_score
```

```
accuracy_scores = []
```

```
running_corrects = 0
```

```
outputs = None
```

```
for inputs, labels in dataloaders['test']:
    model_ft.eval()
```

```
    # print(labels.numpy())
```

```
    inputs = inputs.to(device)
```

```
    labels = labels.to(device)
```

```
    outputs = model_ft(inputs)
```

```
    predicted = torch.max(outputs.data, 1)[1]
```

```
    running_correct = (predicted == labels).sum()
```

```
    running_corrects += running_correct
```

```
accuracy = running_corrects / dataset_sizes['test']
```

```
print("Accuracy: " + str(accuracy.item()))
```

```
Accuracy: 0.625
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 1s completed at 4:04 PM

